

# Deep Learning Lab 2018

**Prof. Wolfram Burgard, Prof. Thomas Brox, Prof. Frank Hutter,  
Asst. Prof. Joschka Boedecker**

University of Freiburg



October 15, 2018

## Exercise 1: Feed-forward Neural Network

- ▶ Submit this exercise done individually until the 30.10.2018 via mail to [kleinaa@cs.uni-freiburg.de](mailto:kleinaa@cs.uni-freiburg.de) and [eitel@cs.uni-freiburg.de](mailto:eitel@cs.uni-freiburg.de).
- ▶ Your submission should include your code (github/bitbucket link) and a 1-2 page report as pdf.
- ▶ **Implementation:** Implement a feed-forward neural network by completing the provided *stub* this includes:
  - ▶ possibility to use 2-4 layers
  - ▶ sigmoid/tanh and ReLU for the hidden layer
  - ▶ softmax output layer
  - ▶ optimization via gradient descent (gd)
  - ▶ optimization via stochastic gradient descent (sgd)
  - ▶ weight initialization with random noise (!!!) (use normal distribution with changing std. deviation for now)
- ▶ Bonus points for testing some advanced ideas:
  - ▶ implement dropout, l2 regularization
  - ▶ implement a different optimization scheme (rprop, rmsprop, adagrad, adam)
- ▶ **Code stub:** <https://github.com/aisrobots/dl-lab-2018>
- ▶ **Evaluation:**
  - ▶ Find good parameters (learning rate, number of iterations etc.) using a validation set (usually take the last 10k examples from the training set)
  - ▶ After optimizing parameters run on the full dataset and test once on the test-set (you should be able to reach  $\approx 1.6 - 1.8\%$  error if you invest some time ;)). Plot the training and validation loss.